

TIME SERIES ANALYSIS

Chirayu Sariya

TIME SERIES ANALYSIS

LEARNING OBJECTIVES

- ▶ Understand what time series data is and what is unique about it
- ▶ Perform time series analysis in Pandas including rolling mean/median and autocorrelation

COURSE

PRE-WORK

PRE-WORK REVIEW

- ▶ Load data with Pandas
- ▶ Plotting data with Seaborn
- ▶ Understand correlation

OPENING

TIME SERIES ANALYSIS

TIME SERIES ANALYSIS

- In this class, we'll discuss analyzing data that is changing over time.
- In most of our previous examples, we didn't care which data points were collected earlier or later than others.
- ▶ We made assumptions that the data was *not* changing over time.
- This class will focus on statistics around data that is changing over time and how to measure that change.

TIME SERIES ANALYSIS

- In this lesson, we will focus on Identifying problems related to time series.
- Additionally, we will discuss the unique aspects of Mining and Refining time series data.

INTRODUCTION

- Time series data is any data where the individual data points change over time.
- This is fairly common in sales and other business cases where data would likely change according to seasons and trends.
- Time series data is also useful for studying social phenomena. For instance, there is statistically more crime in the summer, which is a seasonal trend.

- Most datasets are likely to have an important time component, but typically we assume that it's fairly minimal.
- For example, if we were analyzing salaries in an industry, it's clear that salaries shift over time and vary with the economic period.
- ▶ However, if we are examining the problem on a smaller scale (e.g. 3-5 years), the effect of time on salaries is much smaller than other factors, like industry or position.

- ▶ When the time component *is* important, we need to focus on identifying the aspects of the data that are influenced by time and those that aren't.
- ▶ Typically, time series data will be a sequence of values. We will be interested in studying the changes to this series and how related individual values are.
- For example, how much does this week's sales affect next week's? How much does today's stock price affect tomorrow's?

ACTIVITY: KNOWLEDGE CHECK

ANSWER THE FOLLOWING QUESTIONS



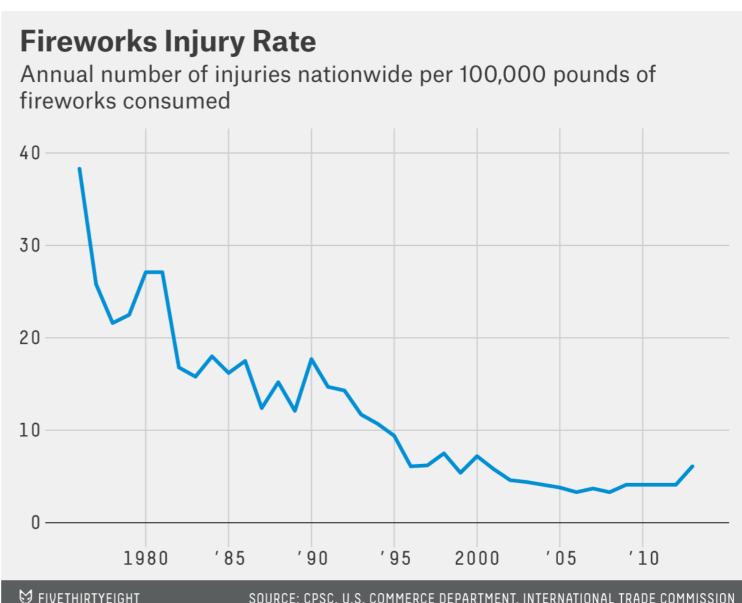
1. Think about the various datasets we've used so far. For each dataset, identify the time components of those datasets. What time related features might be important to our analysis?

DELIVERABLE

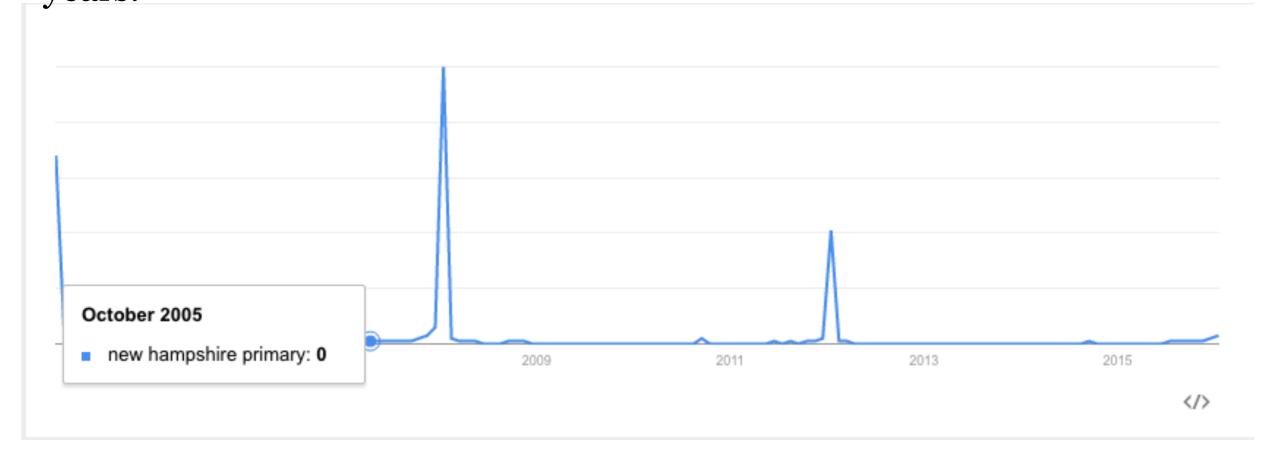
Answers to the above questions

- Time series analysis is useful in many fields: sales analysis, stock market trends, studying economic phenomena, social science problems, etc.
- ▶ Typically, we are interested in separating the effects of time into two components:
 - ▶ Trends significant increases or decreases over time
 - ▶ Seasonality regularly repeating increases or decreases

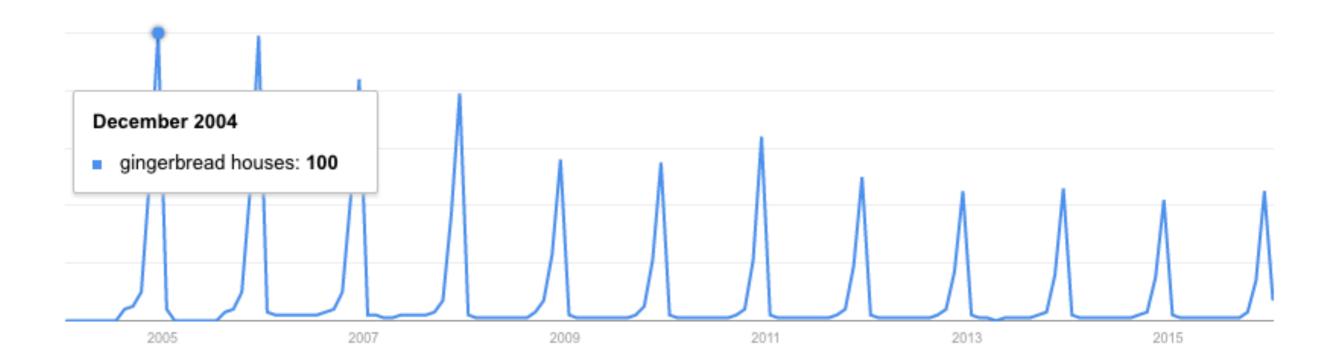
▶ This plot of fireworks injury rates has an overall trend of fewer injuries with no seasonal pattern.



▶ Meanwhile, the number of searches for the New Hampshire Primary has a clear *seasonal* component - it peaks every four years and on election years.

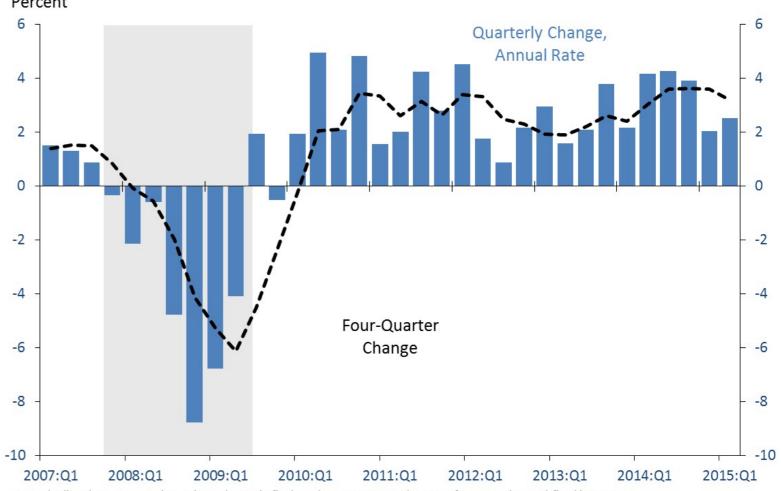


▶ Similarly, searches for 'gingerbread houses' spike every year around the holiday season.



Many other types of regularly occurring up or down swings may occur without a fixed timescale or *period* (e.g. growth vs. recession for economic trends).

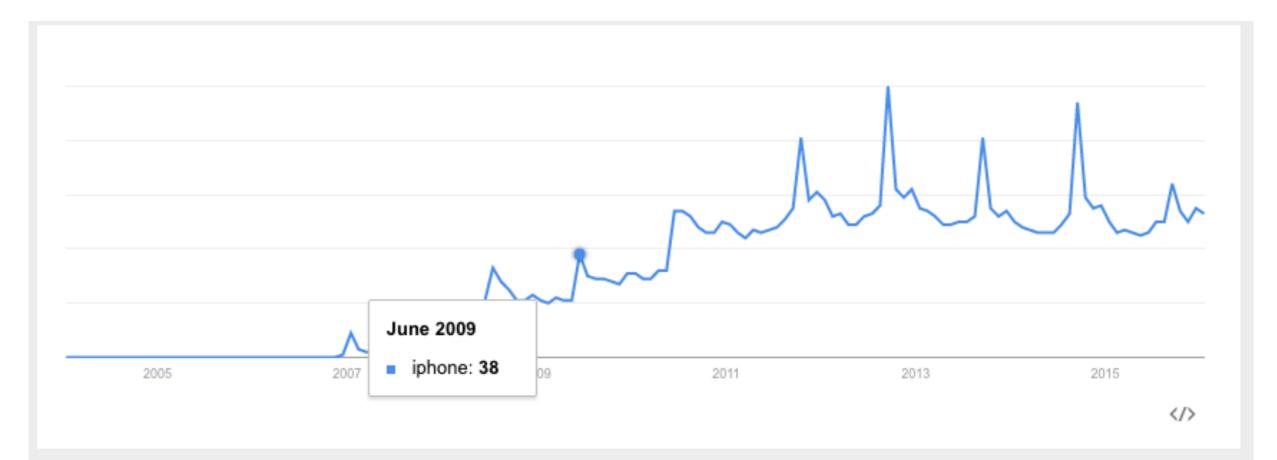
Real Private Domestic Final Purchases Growth, 2007-2015



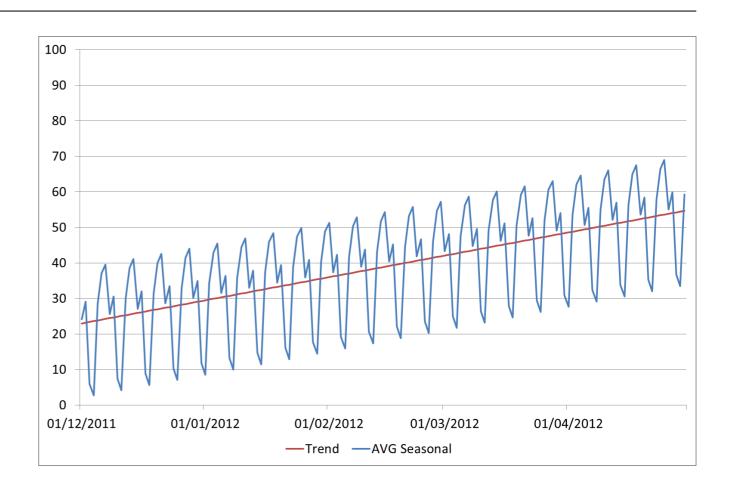
Note: Shading denotes recession. Private domestic final purchases represent the sum of consumption and fixed investment. Source: Bureau of Economic Analysis; CEA calculations.

- ▶ These aperiodic patterns are called *cycles*.
- ▶ While identifying aperiodic cycles is important, they are often treated differently than seasonal effects. Seasonal effects are useful for their consistency, since prior data is useful as a predictor.

Searches for "iphone" have both a general trend upwards (indicating more popularity for the phone) as well as a seasonal spike in September (which is when Apple typically announces new versions).



- Most often, we're interested in studying the *trend* and not the *seasonal* fluctuations.
- Therefore it is important to identify whether we think a change is due to an ongoing trend or seasonal change.



COMMON ANALYSIS FOR TIME SERIES DATA

MOVING AVERAGES

A moving average (or rolling mean) is obtained by taking the average of the initial fixed subset of the number series. Then the subset is modified by "shifting forward"; that is, excluding the first number of the series and including the next value in the subset.



ACTIVITY: KNOWLEDGE CHECK

ANSWER THE FOLLOWING QUESTIONS



1. What would a moving (rolling) mean indicate vs. a moving (rolling) median?

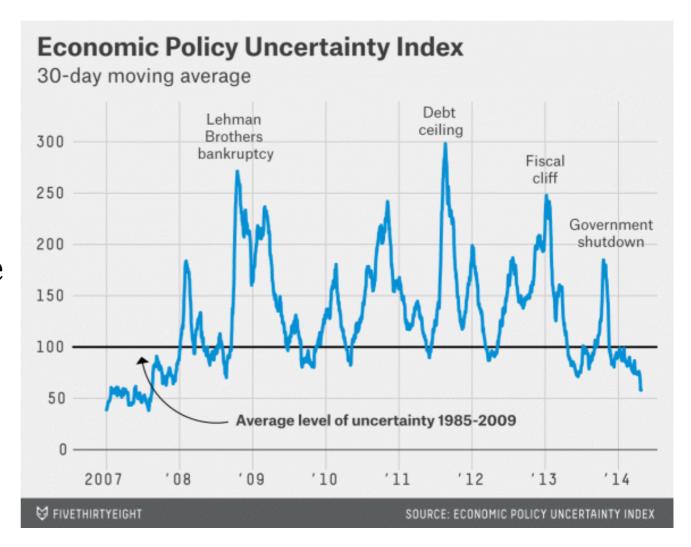
DELIVERABLE

Answer to the above question

MOVING AVERAGES

- A rolling mean would average all values in the window, but can be skewed by outliers (extremely small or large values).
- This may be useful if we are looking to identify atypical periods or we want to evaluate these odd periods.
- For example, this would be useful if we are trying to identify particularly successful or unsuccessful sales days.
- The rolling median would provide the 50 percentile value for the period and would possibly be more representative of a "typical" day.

- This plot shows the 30-day moving average of the Economic Uncertainty Index.
- ▶ Plotting the moving average allows us to more easily visualize trends by smoothing out random fluctuations and removing outliers.

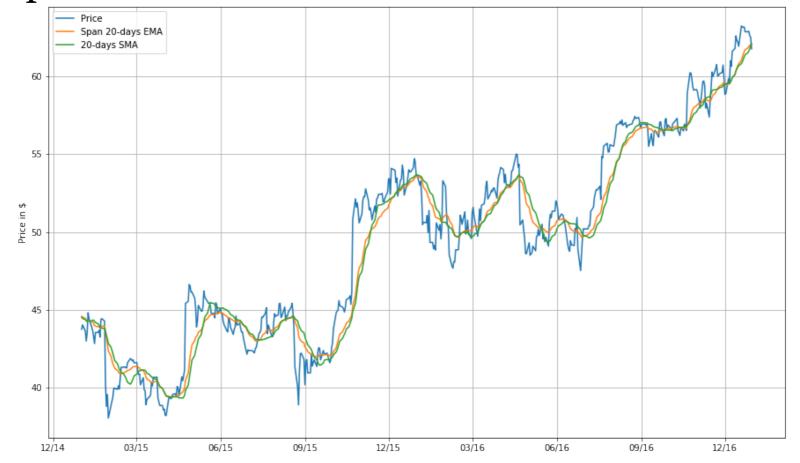


MOVING AVERAGES

- ▶ While this statistic weights all data evenly, it may make sense to weight data closer to our date of interest higher.
- ▶ We do this by taking a *weighted moving average*, where we assign particular weights to certain time points.
- ▶ Various formulas or schemes can be used to weight the data points.

MOVING AVERAGES

A common weighting scheme is an *exponential weighted moving* average (EWMA) where we add a *decay* term to give less and less weight to older data points.



AUTOCORRELATION

- In previous classes, we have been concerned with how two variables are correlated (e.g. height and weight, education and salary).
- *Autocorrelation is how correlated a variable is with itself. Specifically, how related are variables earlier in time with variables later in time.

AUTOCORRELATION

- ▶ To compute autocorrelation, we fix a "lag" *k*. This is how many time points earlier we should use to compute the correlation.
- A lag of 1 computes how correlated a value is with the prior one. A lag of 10 computes how correlated a value is with one 10 time points earlier.

EXPLORING ROSSMANN DRUGSTORE SALES DATA

EXPLORING ROSSMANN DRUGSTORE SALES DATA

- ▶ We will be using data made available by a German drugstore, Rossmann.
- This data contains the daily sales made at the drugstore as well as whether there was a sale or holiday affecting the sales data.
- ▶ Follow along using the starter code available in the class repo.

As always, use Pandas to load our data.

```
import pandas as pd

data = pd.read_csv('../../assets/dataset/rossmann.csv',
    skipinitialspace=True, low_memory=False)
```

▶ Because we are most interested in the Date column, we can process it as a DateTime type and set it as the index of our dataframe.

```
data['Date'] = pd.to_datetime(data['Date'])
data.set_index('Date', inplace=True)

data['Year'] = data.index.year
data['Month'] = data.index.month
```

▶ This allows us to easily filter by date. For example, to a particular year:

```
data['2014']
```

▶ We can also filter to a particular month:

```
data['2015-05']
```

There are over a million sales data points in this dataset, so for some analysis we will focus on just one store.

```
store1_data = data[data.Store == 1]
```

PLOTTING THE SALES DATA

- As we begin to study the sales from this drug store, we will also want to know both the time dependent elements of sales as wells as whether promotions or holidays affected sales.
- ▶ To start, we can simply compare the average sales on those events.
- To compare sales on holidays, we can compare the sales using box plots. This allows us to compare the distribution of sales on holidays against all other days.

- •On state holidays the store is closed (so there should be o sales).
- •On school holidays, the sales are relatively similar.

```
sb.factorplot(
    col='Open',
    x='SchoolHoliday',
    y='Sales',
    data=store1_data,
    kind='box'
)
```

ACTIVITY: KNOWLEDGE CHECK

ANSWER THE FOLLOWING QUESTIONS



1. Check if there is a difference affecting sales on promotion days.

DELIVERABLE

Plot answering the question

• We can see there *is* a difference in sales on promotion days.

```
sb.factorplot(
    col='Open',
    x='Promo',
    y='Sales',
    data=storel_data,
    kind='box'
)
```

▶ Why is it important to separate out days where the store is closed?

- Because there aren't any promotions on those days either, so including them will bias your sales data on days without promotions.
- Let's compare sales across days of the week.

```
sb.factorplot(
    col='Open',
    x='DayOfWeek',
    y='Sales',
    data=storel_data,
    kind='box',
)
```

- Lastly, we want to identify larger scale trends in our data.
- ▶ How did sales change from 2014 to 2015? Were any particularly interesting outliers in terms of sales or customer visits?
- To plot the sales over time:

```
# Filter to days store 1 was open
store1_open_data = store1_data[store1_data.Open==1]
store1 open data[['Sales']].plot()
```

▶ To plot customer visits over time over time:

```
store1_open_data[['Customers']].plot()
```

We can see that there are large spikes of sales and customers towards the end of 2013 and 2014 leading into the first quarter of 2014 and 2015.

ACTIVITY: KNOWLEDGE CHECK

ANSWER THE FOLLOWING QUESTIONS



- 1. Use the index filtering to filter to just 2015. Zoom in on changes over time.
- 2. Is it easier to identify the holiday sales bump?

DELIVERABLE

Plot answering the question

To filter to the 2015 data:

```
store1_data_2015 = store1_data['2015']
store1_data_2015[store1_data_2015.0pen==1][['Sales']].plot()
```

COMPUTING AUTOCORRELATION

- To measure how much the sales are correlated with each other, we want to compute the *autocorrelation* of the "Sales" column.
- In Pandas, we do this with the autocorr function. autocorr takes one argument, lag. This is how many points prior should be used to compute the correlation.

```
data['Sales'].resample('D', how='mean').autocorr(lag=1)
```

As with correlation between different variables, as this number moves closer to 1, the data is more correlated.

ACTIVITY: KNOWLEDGE CHECK

ANSWER THE FOLLOWING QUESTIONS



1. What do the autocorrelation values of "Sales" and "Customers" imply about our data?

DELIVERABLE

Answer to the above question

- If we want to investigate trends over time in sales, we will start by computing simple aggregations. What were the mean and median sales in each year and each month?
- In Pandas, this is performed using the resample function, which is very similar to the groupby function. It allows us to group over different time periods.

- We can use data.resample and provide the following arguments.
 - A level on which to roll up to: 'D' for day, 'W' for week, 'M' for month, 'A' for year.
 - ▶ The aggregation to perform: 'mean', 'median', 'sum', etc.

```
data[['Sales']].resample('A', how=['median', 'mean'])
data[['Sales']].resample('M', how=['median', 'mean'])
```

Here we see that December 2013 and 2014 were the highest average sales months.

- ▶ While identifying the monthly averages are useful, we often want to compare the sales data of a date to a smaller window.
- ▶ To understand holidays' sales, we want to compare the sales data of late December to a few days surrounding it.
- ▶ We can do this using rolling averages.

In Pandas, we can compute the rolling average using the pd.rolling mean or pd.rolling median functions.

This computes a rolling mean of sales using the sales on each day, the day preceding, and the day following (window=3 and center=True).

- rolling_mean (as well as rolling_median) takes the series to aggregate and three important parameters:
 - window the number of days to include in the average
 - center whether the window should be centered on the date or use data prior to that date
 - freq what level to roll up the averages to (as used in resample); 'D' for day, 'W' for week, 'M' for month, 'A' for year

▶ We can use our index filtering to just look at 2015:

```
pd.rolling_mean(data[['Sales']], window=3, center=True, freq='D')['2015']
```

Instead of plotting the full time series, we can plot the rolling mean instead. This smooths random changes in sales as well as removing outliers, helping us identify larger trends.

```
pd.rolling_mean(data[['Sales']], window=10, center=True, freq='D').plot()
```

- As we discussed earlier, this averages all values in the window evenly. However we may want to weight closer values more.
- ▶ For example, for a centered weighted average of 10 days, we want to put emphasis on +/- 1 day versus +/- 5 days.
- •One option to do that is the ewma function or the exponential weighted moving average function.

```
pd.ewma(data['Sales'], span=10)
```

ACTIVITY: KNOWLEDGE CHECK

ANSWER THE FOLLOWING QUESTIONS



1. Discuss the differences between the plot of the 10-day moving average and the 10-day exponentially weighted moving average.

DELIVERABLE

Discussion

PANDAS WINDOW FUNCTIONS

- ▶ Pandas rolling_mean and rolling_median are only two examples of Pandas window function capabilities.
- ▶ Window functions operate on a set of N consecutive rows (a window) and produce output.
- In addition, there are rolling_sum, rolling_min, rolling_max, and many more.

PANDAS WINDOW FUNCTIONS

- ▶ Another common window function is diff, which takes the difference over time.
- pd.diff takes one argument, periods, which is how many rows prior to use for the difference.
- ▶ For example, if we want to compute the difference in sales, day by day:

```
data['Sales'].diff(periods=1)
```

PANDAS WINDOW FUNCTIONS

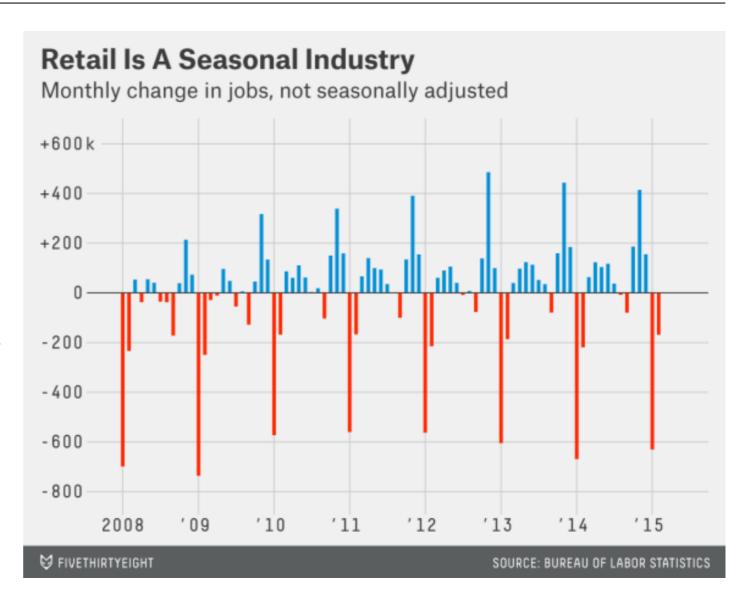
▶ However, if we wanted to compare the same day in the prior week, we could set periods=7.

```
data['Sales'].diff(periods=7)
```

This would compute the difference in sales, from every day to the same day in the previous week.

WHAT IS TIME SERIES DATA?

- Difference functions allow us to identify seasonal changes as we see repeated up or down swings.
- This plot of the month to month change (diff) in retail jobs helps identify the seasonal component of the number of retail jobs.



PANDAS EXPANDING FUNCTIONS

- In addition to the set of "rolling" functions, Pandas also provides a similar set of "expanding" functions.
- Instead of using a window of N values, "expanding" functions use all values up until that time.

PANDAS EXPANDING FUNCTIONS

▶ We can compute the average sales from the first date *until* the date specified.

```
pd.expanding_mean(data['Sales'], freq='d')
```

▶ We can also compute the sum of average sales per store up until that date.

```
pd.expanding_sum(data['Sales'], freq='d')
```

INDEPENDENT PRACTICE

TIME SERIES EXERCISES

ACTIVITY: TIME SERIES EXERCISES



DIRECTIONS (35 minutes)

- 1. Plot the distribution of sales by month and compare the effect of promotions.
- 2. Are sales more correlated with the prior date, a similar date last year, or a similar date last month?
- 3. Plot the 15 day rolling mean of customers in the stores.
- 4. Identify the date with largest drop in sales from the same date in the previous month.
- 5. Compute the total sales up until Dec. 2014.
- 6. When were the largest differences between 15-day moving/rolling averages? **HINT**: Using rolling_mean and diff

DELIVERABLE

Plots and answers to the above questions

INDEPENDENT PRACTICE REVIEW

TIME SERIES EXERCISES

▶ Plot the distribution of sales by month and compare the effect of promotions.

```
sb.factorplot(
    col='Open',
    hue='Promo',
    x='Month',
    y='Sales',
    data=storel_data,
    kind='box'
)
```

Are sales more correlated with the prior date, a similar date last year, or a similar date last month?

```
# Compare the following:
average_daily_sales = data[['Sales', 'Open']].resample('D', how='mean')
average_daily_sales['Sales'].autocorr(lag=1)
average_daily_sales['Sales'].autocorr(lag=30)
average_daily_sales['Sales'].autocorr(lag=365)
```

▶ Plot the 15 day rolling mean of customers in the stores.

```
pd.rolling mean(data[['Sales']], window=15, freq='D').plot()
```

▶ Identify the date with largest drop in sales from the same date in the previous month.

```
average_daily_sales = data[['Sales', 'Open']].resample('D', how='mean')
average_daily_sales['DiffVsLastWeek'] =
average_daily_sales[['Sales']].diff(periods=7)
average_daily_sales.sort(['DiffVsLastWeek']).head
```

▶ Unsurprisingly this day is December 25th and 26th in 2014 and 2015, when the store is closed and there were many sales in the preceding week.

Compute the total sales up until Dec. 2014.

```
total_daily_sales = data[['Sales']].resample('D', how='sum')
pd.expanding_sum(total_daily_sales)['2014-12']
```

Note that this is **NOT**

```
pd.expanding_sum(data['Sales'], freq='D')
```

▶ We don't want to average over stores first!

When were the largest differences between 15-day moving/rolling averages? **HINT**: Using rolling_mean and diff

```
pd.rolling_mean(data[['Sales']], window=15,
freq='D').diff(1).sort('Sales')
```

▶ Unsurprisingly, they occur at the beginning of every year after the holiday season.

CONCLUSION

TOPIC REVIEW

CONCLUSION

- ▶ We use time series analysis to identify changes in values over time.
- ▶ We want to identify whether changes are true trends or seasonal changes.
- Rolling means give us a local statistic of an average in time, smoothing out random fluctuations and removing outliers.
- Autocorrelations are a measure of how much a data point is dependent on previous data points.

LESSON

Q&A

LESSON

EXIT TICKET

DON'T FORGET TO FILL OUT YOUR EXIT TICKET